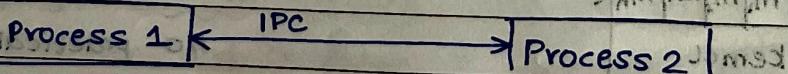


## Module-3

- \* IPC (Inter-process Communication)
- 1. Independent Process:
  - Process cannot affect or cannot be affected by other processes
- 2. Cooperating Process:
  - can affect or be affected by other processes
  - Any process that shares data with other process
  - Reasons:
    - (i) Several users might be interested in the same piece of information. We must provide an environment to allow concurrent access to such information. (Information Sharing)
    - (ii) If we want a task to run faster, we must break it into sub tasks and each subtask will execute in parallel. (Computation Speedup)
    - (iii) Speedup can be achieved if computer has multiple processing elements like CPUs or I/O channels.
    - (iv) We must divide the system functions into processes or threads. (Modularity)
    - (v) Individuals can also work on many tasks at the same time (Convenience)
  - Cooperating processes need IPC to exchange data and information.



versus 20

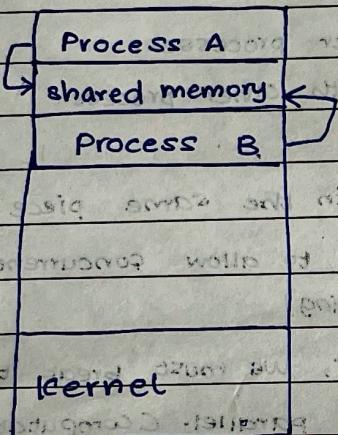
## \* IPC

- Two models of IPC:

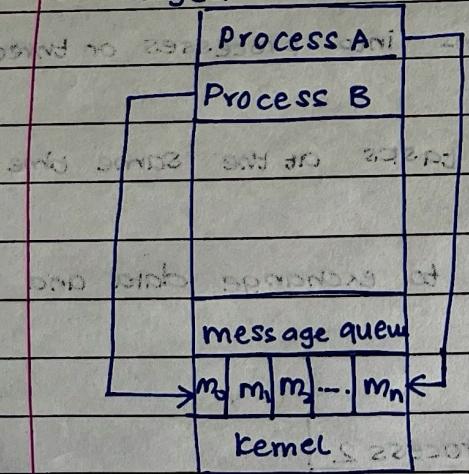
. Shared Memory

. Message Passing

A region of memory is shared by cooperating process is established. Processes can read and write data into the shared region.



Communication takes place by means of messages that are exchanged between cooperation of processes.



\* Unbounded buffer

no practical limit on the size of buffer.

## Message Passing vs Shared Memory

- Message passing is useful when smaller amounts of data are exchanged because no conflicts can be avoided.
- Slower as message passing systems are typically implemented using system calls and thus the required only to establish more time-consuming task of shared regions. Once the kernel intervention. Shared regions are established, all accesses are treated as memory access and no assistance from kernel is needed.

### → Shared Memory

#### \* Producer - Consumer Problem (Bounded Buffer Problem)

- Paradigm for cooperating processes: Producer process produces information that is consumed by consumer process.
- eg: Compiler produces assembly code which is consumed by an assembler, the assembler produces object modules which are consumed by loader.
- eg: Server (Web Server) produces HTML files and images which are consumed by Web Browser requesting the resource.

#### \* Solution 1: Shared Memory Problem Statement

- Allow Producer and consumer to run concurrently.
- Available buffer of items that can be filled by producer and emptied by consumer.
- Buffer resides in memory that will be shared by producer & consumer processes.

S	M	T	W	T	F
AVIUDY					

M	T	W	T	F
Page No.:				

Page No.:

Date:

Consumer  
next item  
the  
(Solution)

- Producer is producing one item, while consumer consumes

two items

Producer ↓

available space in buffer

Consumer ↓

produced before in buffer

SYN - Producer and consumer processes must be synchronised so that producer does not produce when buffer is full and consumer doesn't consume when buffer is empty.

- If an unbounded buffer was used here then, consumer would have to wait for the new items to be produced and producer can always produce new items.

If a bounded buffer was used, then consumer must wait if buffer is empty and producer must wait if buffer is full.

- Algo:

Following variables reside in a region shared by memory shared by producer and consumer process:

Shared buffer is implemented as a circular array with two pointers:

in - points to next free position in buffer

out - points to first full position in buffer

#define BUFFER\_SIZE 10

typedef struct {

} item;

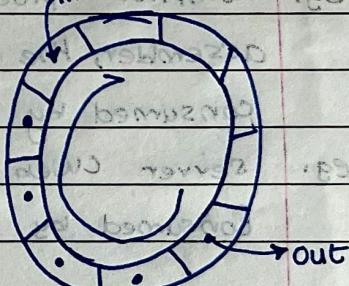
item buffer[BUFFER\_SIZE];

int in, out;

Buffer is empty if in == out

Buffer is full if when  $((in+1) \% \text{BUFFER\_SIZE}) == \text{out}$

Producer process has a new variable nextProduced in which the new item to be produced is stored.





M	T	W	T	F	S	S

M	T	W	T	F	S	S
Page No.:						

Date: YOUVA

- Consumer Process has new local variable `nextConsumed` that stores the next item to be consumed.
- (Solution - 2 done later)

### 220 \* Message Passing

- Message Passing is another way of cooperating processes to communicate with each other.
- Mechanism for processes:
  - (i) to communicate
  - (ii) to synchronize their actions.
- Processes communicate with each other without resorting to shared variables or shared address space.

Used by a distributed system environment where communicating processes reside on different computers. They are connected by a network.

Eg: Chat application

IPC provides two operations: `sendC message`, `receiveC message`

Message size is fixed or variable.

If fixed-size message is sent, implementation is straightforward and programming task is difficult.

If variable-size message is sent, implementation is complex but programming task is easy.

If two processes, A and B wish to communicate, a communication link must be made between them and messages can be exchanged using `send()` or `receive()` functions.

Two ways of communication between processes:

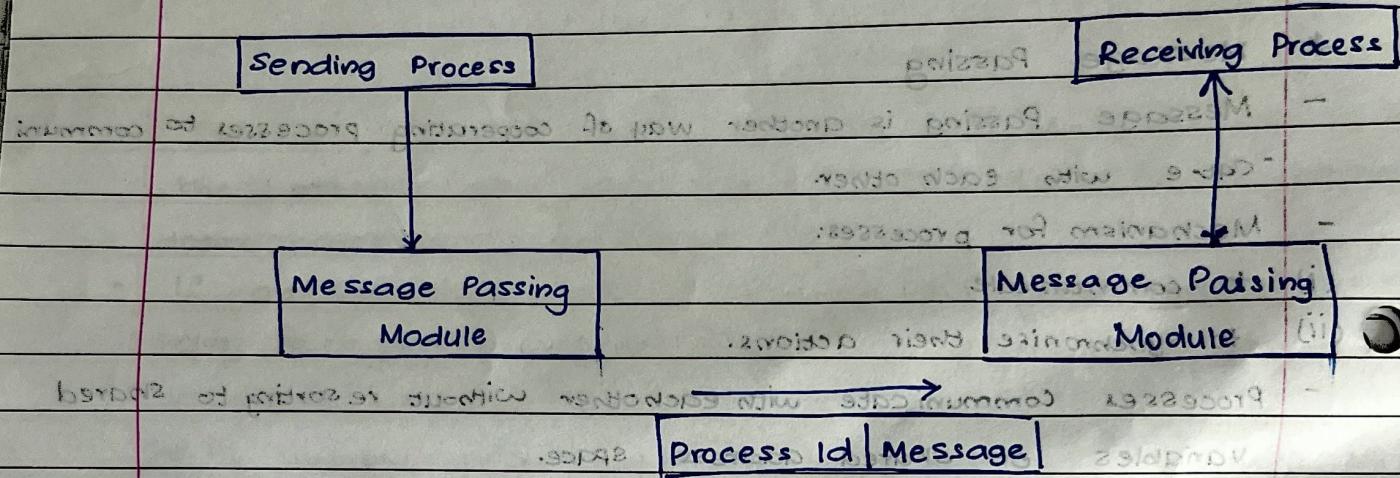
- (i) Direct Communication
- (ii) Indirect Communication

Sender Receiver

S	R	T	W	T	M
A	V	U	O	N	E
Paper No.:		Date:			

### \* Direct Communication

- Processes must name each other explicitly.
- . send (P, message): send a message to P.
- . receive (Q, message) receive a message from Q.



Processes must know only one another's identity when they wish

to communicate.

- A link is associated with only one pair of communicating process.

- Link can be bidirectional or unidirectional.

Schemes exhibit symmetry, i.e., must name each other to communicate.

Assymetry is employed when the sender names the recipient and recipient is not required to name the sender.

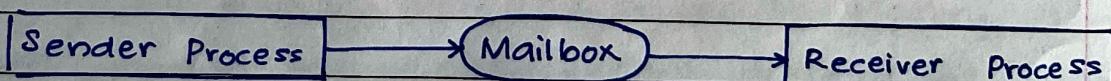
### \* Indirect Communication:

Messages are sent and received from mailboxes (ports).

Each mailbox has a unique id.

Processes communicate only if they share a mailbox.

Mail box can be viewed as an object in which messages are placed by processes and also removed.



S	T	W	T	M
Page No.:				
Date:				

M	T	W	T	F	S	S
Page No.:						YOUVA

- One process can communicate with other processes via a number of different mailboxes.
- Link is established only if processes share a common mailbox.
- Each link corresponds to only one mailbox.
- Links can be bidirectional or unidirectional.
- send (a, message) - send a message to mailbox 'a'.
- receive (b, message) - send a message to mailbox 'b'.
- Mailbox is owned either by a process or operating system.
- If mailbox is owned by a process, then it becomes a part of the address space of the process.
- The owner can only receive messages from the mailbox.
- The user can only send messages to the mailbox.
- Since each mailbox has a unique owner, there is no confusion about which process sent to the mailbox.
- If the process that owns the mailbox terminates then the mailbox also disappears and all the processes that send a message to this mailbox are also notified.
- If mailbox is owned by OS, then mailbox has an existence of its own, it is independent and not attached to any process.
- OS then allows the process to:
  - create a new mailbox
  - send and receive messages through mailbox
  - delete the mailbox.
- Process becomes owner of mailbox.
- Ownership and receiving privileges can be passed using system calls. This will cause multiple receivers for each mailbox.

S	M	T	W	T	F	S
A	V	U	O	D	Y	E
1	2	3	4	5	6	7
8	9	10	11	12	13	14

M	T	W	T	F
Page No.:	100	Date:	10/10/2023	Y05

### \* Synchronization

- IPC takes place via send() or receive() primitives.

X - Message passing can be done in two ways:

- (i) Blocking send (Synchronous)

- Sending process is blocked until the receiving process or the mailbox receives the message.

- (ii) Blocking receive

- Receiver is blocked until message is available.

- (iii) Non-blocking send (Asynchronous)

- Sender process sends message and continues with operation.

- (iv) Non-blocking receive (Asynchronous)

- Receiver receives a valid message or null.

→ Solution - Producer and consumer both implement blocking send and receive.

→ Producer-consumer problem can be solved using synchronization.

- Producer invokes blocking send() call and waits till message is delivered to receiver or mailbox.

- Consumer invokes blocking receive() until a message is available.

### \* Buffering

- Whether message communication is direct or indirect, messages are exchanged and they reside in a temporary queue.

- Queues are implemented in three ways:

#### ① Zero capacity:

- Max. Queue Length = 0
- Link cannot have any message waiting in it.
- The sender must block until the recipient arrives.

#### ② Bounded Capacity:

- Queue has finite length 'n'.
- Almost 'n' messages can reside in it.
- Link's capacity is infinite.
- If queue is not full, message is placed in queue. Either the message is copied or a pointer to the message is kept.
- If queue is full, sender must block until space is available.

Unbound  
Queue

M	T	W	T	F	S
Page No.:	200	Left Margin:	9869	Right Margin:	1000
Date:	10/10/2023				

M	T	W	T	F	S
Page No.:	200	Left Margin:	9869	Right Margin:	1000
Date:	10/10/2023				

### (3) Unbounded Capacity:

- Queue length is potentially infinite.
- Any no. of messages can wait in it.
- Sender never blocks.
- zero capacity with no buffering.

### \* Producer - Consumer Problem

#### Solution-1

- Modifying the approach for Producer- Consumer Problem by having a counter that keeps track of total number of buffers (full).
- One possibility is to have an integer variable counter, initialized to 0. It is incremented every time we add a new item to buffer and decremented when we remove one item from the buffer.
- Earlier Approach :

```

    item = nextProduced;
    while (true) {
        while ((Cinti) % BUFFER_SIZE == out);
        buffer [in] = nextProduced;
        in = (in + 1) % BUFFER_SIZE;
        set buffer counter +1;
    }

```

PRODUCER

#### - Modified Approach :

```

    while (true) {
        while (counter == BUFFER_SIZE);
        /* do nothing */
        buffer [in] = nextProduced;
        in = (in + 1) % BUFFER_SIZE;
        Counter++;
    }

```

S	M	T	W	T	F
MONDAY					
Date:					

M	T	W	T	F
Page No.:				

- Earlier Approach (Consumer Process)

```

item nextConsumed;
while (true) {
    while (in == out),
        /* do nothing */;
    nextConsumed == buffer [out];
    out = (out + 1) % BUFFER_SIZE;
}

```

- Modified Approach:

```

while (true) {
    while (counter == 0)
        /* do nothing */;
    nextConsumed == buffer [out];
    out = (out + 1) % BUFFER_SIZE;
    counter--;
}

```

- Producer and consumer processes work correctly separately but may not execute correctly when executed concurrently.
- When several processes access and manipulate data concurrently, outcome depends on order in which access takes place. It is also called Race condition.
- To guard against race condition, only one process must manipulate the variable counter. Thus, process should be synchronised.

(Semaphore Solution done later)

$$\begin{aligned}
 & \text{initially } \text{buffer} = [0] \\
 & \text{buffer}[0] = \text{buffer}[0] + 1 \\
 & \text{buffer}[0] = 1 + (\text{buffer}[0] - 1)
 \end{aligned}$$

### \* Critical Section Problem:

- Each process has a segment of code, called critical section in which the process may be changing common variables, writing a file, updating a table and so on.
- When one process is executing in its critical section, no other process is allowed to execute in its critical section. Execution of critical sections is mutually exclusive.

Entry Section: Each code must request a permission to enter its critical section. This section of implementing the request is called entry section.

- Exit Section: The critical section is followed by an exit section.
- Remaining Section: The remaining code is the remaining section.
- Critical Section Solution: Solution must satisfy three requirements:

(i) Mutual Exclusion: If a process  $P_i$  is executing in their critical section, no other process can be executing in their critical sections.

(ii) Progress: If no process is executing in critical section and some processes wish to enter the critical section, then only those processes can enter which are not executing in remainder section can participate in deciding which will enter the critical section next.

This selection cannot be postponed indefinitely. Only those process interested in entering CS should compete to enter the CS.

(iii) Bounded Waiting: There exists a limit / bound on the no. of times that the processes are allowed to enter critical sections. Max after a bound / time limit, after which the process definitely will get a chance to enter critical section.

## \* Critical Section Handling in OS

- There are 2 approaches to handle critical sections in OS depending on if kernel is:
  - (i) Pre-emptive :
  - Preemptive kernels allow a process to be pre-empted while it is running in kernel mode.
  - Two kernel mode processes can run simultaneously on different processors.
  - Must be carefully designed to ensure kernel data are free from race conditions.

### (ii) Non-preemptive:

- Does not allow kernel in a process to be preempted.
  - Free from race conditions on kernel data structures.
  - Only one process is active in kernel, one at a time.
- Non-preemptive Kernel Approach will run the process until it:
- exits kernel mode,
  - blocks or
  - voluntarily yields control of CPU.
- A preemptive kernel is more suitable for real-time programming, as it will allow a real-time process to preempt a process currently running in the kernel. It may be more responsive, since there is less risk that a kernel-mode process will run for an arbitrarily long period before relinquishing the processor to waiting processes.

M	T	W	T	F	S	S
Page No.:						YOUVA
Date:						

## \* Solutions to Critical Section Problem

### 1. Software based solutions:

#### → Two Process Solutions:

- Assume that the load and store machine-language instructions are atomic - that is, cannot be interrupted.

#### (A) Algorithm 1 :

do {

while (turn != i); // entry code

critical section

turn = j; // exit code

remainder section

} while (1);

- Structure of Process P<sub>i</sub> in algorithm 1: i = 0 and 1.
- Entry code acts like a trap, it stops the processes from entering in the critical section.
- Turn = Shared, Common / Global Integer initialised to 0 or 1.
- Checking for Mutual exclusion:

• Initialize turn with 0.

P<sub>0</sub>, i=0

P<sub>1</sub>, i=1

do {

processes { op

while (turn != 0);

while (turn != 1);

critical section

critical section

turn = 1;

turn = 0;

remainder section

} while (1);

} while (1);

Can P<sub>1</sub> enter CS if P<sub>0</sub> is in CS?

- Lets take turn = 0 for P<sub>0</sub>.
- While condition = false, P<sub>0</sub> enters critical section.
- While P<sub>0</sub> is still in CS, turn = 0.
- P<sub>1</sub> tries to enter CS. Turn is still 0.

- while condition = true,  $P_1$  gets trapped in an infinite loop.
- $P_1$  is unable to enter CS.

- Can  $P_1$  enter CS while  $P_0$  is in remainder section?

- For  $P_0$ , after critical section is executed, the exit code makes turn = 1 which stops both  $P_0$  and  $P_1$ .
- It starts executing remainder section.
- $P_1$  tries to enter CS.
- while condition = false,  $P_1$  enters CS.

- Can  $P_0$  enter CS again after completing RS?

- After first run, turn = 1
- For  $P_0$ , while condition = true,  $P_0$  gets trapped in an infinite loop.
- $P_0$  cannot enter CS.

MUTUAL EXCLUSION SATISFIED!

- if turn == i,  $P_i$  executes in CS
- Progress Requirement Checked
- After  $P_0$  comes out,  $P_1$  gets to go inside and vice-versa.
- If a process does not want to go in CS, it still has to go due to alteration.
- If  $P_1$  does not want to go and  $P_0$  wants to go in CS, it still cannot go.

PROGRESS REQUIREMENT NOT SATISFIED!

- Bounded Waiting Check:
- Process  $P_0$  can go directly and Process  $P_1$  can go immediately after  $P_0$  is done with the critical section.

• BOUNDED WAITING SATISFIED!

- Mutual Exclusion is Preserved: Ensures that only one process at a time can be in its critical condition.
- Progress Requirements not satisfied.
- Bounded Waiting satisfied.

### (B) Algorithm 2:

Two Processes share a Boolean array.

- The flag array is used to indicate if a process is ready to enter the critical section.  $\text{flag}[i] = \text{true}$  implies process  $P_i$  is ready.
- Boolean array is initialised to be false.
- As it wants to enter CS, it is changed to True.

$\text{do} \{$       MUTUAL EXCLUSION GUARANTEE

$\text{flag}[i] = \text{true};$

$\text{while } (\text{flag}[j]);$

$\text{flag}[i] = \text{false};$

remainder section

$\} \text{ while}(i);$

Structure of process in Algorithm 2

$P_0$

$\text{do} \{$

$\text{flag}[0] = \text{true};$

$\text{while } (\text{flag}[1]);$       while  $(\text{flag}[0]);$

critical section

$\text{flag}[1] = \text{false};$        $\text{flag}[0] = \text{false};$

remainder section

remainder section

$\} \text{ while}(0);$        $\} \text{ while}(1);$

- Checking for Mutual Exclusion:

- If  $P_0$  is executing in CS, can  $P_1$  enter CS?

If  $P_0$  wants to enter CS, set  $\text{flag}[0] = \text{T}$ .

-  $P_0$  checks if  $P_1$  wants to go to CS.

- As  $\text{flag}[1] = \text{F}$ , while condition is false,  $P_0$  enters critical section and executes it.

[0]	[1]
T	F

[0]	[1]
T	T

- $P_1$  tries to enter CS.

- Set  $\text{flag}[1] = \text{T}$ .

- While condition = true,  $P_1$  goes in an infinite loop and does not enter CS.

	S	M	T	W	T	F
AVIJOY	9	10	11	12	13	14
Day	1	2	3	4	5	6

M	T	W	T	F
Page No.:				
Date:				

- Can  $P_1$  enter CS if  $P_0$  is in Remainder Section?
  - For  $P_0$ , after executing CS,  $\text{flag}[0]=\text{false}$  and it enters RS.
  - $P_1$  tries to enter CS. Set  $\text{flag}[1]=T$ .
  - While condition is false, so  $P_1$  enters CS.

[0]	[1]
F	T

MUTUAL EXCLUSION SATISFIED!

### Progress Requirements Check:

- Can  $P_0$  enter CS again after completing RS?
  - If  $P_1$  is not interested in CS and  $\text{flag}[1]=F$  and  $P_0$  wants to go again, we set  $\text{flag}[0]=T$ .
  - While condition = false,  $P_0$  enters CS.

[0]	[1]
T	F
F	F

### Progress Requirement Check:

- If  $P_1$  wants to go in CS again and immediately after RS, it can enter if  $P_2$  is not interested.
- Whichever process interested and other processes are not, can enter the CS.
- If both  $P_0$  and  $P_1$  get trapped in while loop, both cannot enter CS.

[0]	[1]
T	T

Both go in an infinite loop and thus there is no progress here.

- Mutual Exclusion preserved.

Progress requirement is not satisfied.

[0]	[1]	[2]	[3]
T	T	1	T

23/06/2023 at 10:19

$T = [0] \oplus [1] \oplus [2]$

A.23/06/2023 10:19 hrs good simulation of 23/06/2023, and = condition given

S	T	W	T	M
YOUVA				

M	T	W	T	F	S	S
Page No.:	YOUVA					
Date:						

### (c) Algorithm 3 / Peterson's Solution

- Two processes share two variables;

[1] int turn; and Boolean flag[2]

Variable turn indicates whose turn it is to enter CS and flag array indicates if a process is ready to enter CS.

do {

[1] [2] flag[i] = True; i=1,2 loop 23

T T turn = j;

while (flag[j] && turn == j);

Critical Section

flag[i] = False;

{ remainder section is in i = 0 to 1

3 while (true); i = 0 to m to execute

Structure of Peterson's Solution

P0

do {

flag[0] = true;

turn = 1;

while (flag[1] == T && turn == 1);

Critical Section

flag[0] = false;

5 while (true);

do {

flag[1] = True;

turn = 0; i = 1 to m to execute

while (flag[0] == T && turn == 0);

Critical Section

flag[1] = False;

5 while (1);

- Mutual Exclusion Check:
  - If  $P_0$  is in CS, can  $P_1$  enter CS?
    - For  $P_0$ , set flag as true and turn as 1.
    - $P_1$  is not interested
    - while ( $T \& F$ )  $\rightarrow$  while condition is false.
    - $P_0$  enters the CS.
    - $P_0$  is already in CS and  $P_1$  tries to enter CS.
    - Set flag  $[1] = T$  and turn = 0
    - while ( $T \& T$ )  $\rightarrow$  while condition is true.  $P_1$  gets trapped in infinite loop. So it can't enter CS.

- Can  $P_1$  enter while  $P_0$  is in remainder section?
  - After  $P_0$  is done with execution is in CS, flag is set as False
  - $P_1$  tries to enter CS. Flag  $[1] = T$  and turn = 0.
  - while ( $T \& F$ )  $\rightarrow$  while condition is false.
  - $P_1$  enters CS

$[0]$	$[1]$
F	T

turn = 0

MUTUAL EXCLUSIVE = SATISFIED!

- Progress Requirement Check:
  - If both  $P_0$  and  $P_1$  want to enter CS
    - For  $P_0$ , if  $P_1$  is interested. Set flag as true and turn == 0
    - Context switch occurs.
    - As  $P_0$  had set turn as 1
    - Now,  $P_1$  tries while ( $F \& T$ )  $\rightarrow$  while condition is false.  $P_1$  enters CS.
    - Now,  $P_0$  is also interested. Set flag as true and turn = 1
    - while ( $T \& T$ )  $\rightarrow$  while condition is true.  $P_0$  gets trapped.
    - Now context switch.  $P_0$  does not enter CS.

PROGRESS CHECK SATISFIED!

### Bounded Waiting Check:

P<sub>i</sub> will be critical section after atmost One entry by P<sub>j</sub>. Everyone gets a fair chance.

BOUNDED WAITING CHECK SATISFIED!

### Hardware based Solutions:

Software solutions like Peterson's may not work for modern computer architectures.

Race conditions are prevented and critical regions can be protected by locks.

A process must acquire a lock before entering CS; it is released when it exits the CS.

Critical section

• Release lock;

remainder section

? while (1);

(A)

Test Memory word and set value: test and set()

boolean test-and-set (boolean \*target)

{

    boolean rv = \*target;

    \*target = TRUE;

    rv = \*target; return rv;

}

Executed automatically. If 2 tests and set instructions are executed simultaneously, they will be executed sequentially in some arbitrary order.

Returns original value of passed parameter and sets new parameter value to True.

## Mutual Exclusion implementation

do {

    while (!TestAndSetC(&lock));

        critical section

    lock = False;

    remainder section

} while();

P<sub>0</sub> tries to enter CS, set lock = true and TBS returns false.

- while (false), P<sub>0</sub> enters CS.

Now, if P<sub>1</sub> tries to enter CS, P<sub>0</sub> is still in CS.

- lock = true and TBS returns true.
- while (true), thus P<sub>1</sub> goes in infinite loop. P<sub>1</sub> does not enter CS.
- After P<sub>0</sub> has come out of CS, lock becomes false and it goes in remainder section.
- Now P<sub>1</sub> tries to enter CS as lock is false.
- while condition is now false and P<sub>1</sub> comes out of loop and enters CS.

## Swap Instruction

void Swap(boolean \*a, boolean \*b)

    boolean temp = \*a;

    \*a = \*b;

    \*b = temp;

}

; swap = TRUE;

See for all in PDF

Basics and Applications of Computer Organization

Switches directly to semaphores now.

Both test and set, swap do not satisfy bounded wait requirement.

## \* Semaphores

- Semaphore is a synchronization tool.
- It is an integer  $S$  that is accessed apart from initialization only through two standard operations:  $\text{Wait}()$  and  $\text{Signal}()$ .
- $\text{Wait}()$  operation originally termed 'P' from proberen.
- $\text{Signal}()$  operation originally termed 'V' from verbogen.
- $\text{Wait}(S)$  {  $S--$  }
- $\text{Signal}(S)$  {  $S++$  }
- When one process is modifying the semaphore value, no other process can simultaneously modify the semaphore value.

## \* Counting semaphore

- The value of a counting semaphore can range over an unrestricted domain.
- The semaphore is initialized to the number of resources available.
- When a process uses a resource, it performs  $\text{wait}()$  operation on the semaphore and decreases the count.
- When a process releases a source, it performs  $\text{signal}()$  operation and increases the count.
- When the count reaches 0, all resources are used. After that if a process wishes to use a resource, it will be blocked until the count becomes greater than 0.

## \* Binary Semaphore

- Value can range only between 0 and 1.
- Also called mutex locks
- Provide mutual exclusion.

### \* Busy Waiting:

- No two processes can execute the `Wait()` and `Signal()` on the same semaphore at the same time.
- When a process tries to enter a CS while another process is already in its CS, it must go in an infinite loop in the entry code.
- Busy waiting wastes CPU cycles that some other process can use productively.
- All software solutions suffered from busy waiting.
- When a process is in busy waiting critical section, the semaphore used is spin lock. The process 'spins' while waiting for lock.
- When locks are to be held for a short time, spin locks are useful.

Note - One thread can spin while other thread is in CS.

(See implementation of semaphore with no busy waiting)

### \* Bounded Buffer Problem

(Producer-Consumer soln - 3 using semaphores)

Semaphore & mutex - Binary semaphore provides mutual exclusion for access to buffer pool.

Initialized value to 1. It provides local storage for one item.

Semaphore full - Counting semaphore counts no of full buffers.

Initialized to 0. It provides local storage for one item.

Semaphore empty - Counting semaphore counts no of empty semaphores.

Initialized to n. It provides local storage for one item.

S	T	W	T	F
Page No.:				
Date:				

M	T	W	T	F	S	S
Page No.:						

## - Structure of Producer Process

do {

; (Consumer) Wait (empty);

; (Producer) wait (mutex);

// CS: add item to buffer

; (Consumer) signal (mutex);

(Producer) signal (full);

} while (1);

; (Consumer) tidw

## \* Reader's - Writer's problem

Readers can only read dataset, do not perform any updates

writers can both read and write or update.

### - Problem:

- Allow multiple readers to read at same time, no adverse effect
- If writer and some other process (reader / writer)

### - Requirement:

- Writer must have exclusive access to shared object.

### 1. Reader's - Writer's Problem:

#### - Semaphore rw-mutex → Binary semaphore

Initialised to 1

Common to both reader and writer

Mutual exclusion for writers

Also used by first or last reader that enters or exits mutex.

#### - Semaphore - mutex → Binary semaphore

Initialised to 1 and releases reading R when reader exits

Ensure mutual exclusion when read\_count is updated.

integer read\_count; if read\_count > 0 then read\_count = read\_count - 1;

Initialised to 0 and increments R if read\_count > 0;

keeps track of how many processes are currently reading the object.

S	E	R	T	W	T	M
A	V	O	Y			
Page No.:						

M T W T F

Page No.:

Date:

YODA

### - Writer Process -

do {

for

{ Wait(rw-mutex);  
// writing is performed

signal(rw-mutex);

3. while (1) :

{ (some) logic

{ do while }

### - Reader Process -

do {

for

to mutex < Wait(mutex);

read-count variable (read-count++;

if(read-count == 1)

ensures no writer  
can enter if there is  
even one reader

{ signal(mutex);

// reading is performed

reader

leaves

{ Wait(mutex);

read-count--;

if(read-count == 0)

writer can enter < signal(rw-mutex);

enter

reader leave < signal(mutex);

### 2. Writer- Reader Problem:

Same set of code due to race condition between writer and reader.

### 3. Writer- writer problem

Take only writer code

### 4. Reader - reader problem

Take only reader problems.

### \* Dining Philosopher's Problem.

- Philosophers spend their life thinking and eating

- When thinking, they do not interact with neighbors

- When eating/hungry, a philosopher needs both left and right

chopstick to eat.

- A hungry philosopher will not eat if both chopsticks are not available

- A philosopher will pick up 2 chopsticks closest to her to eat.

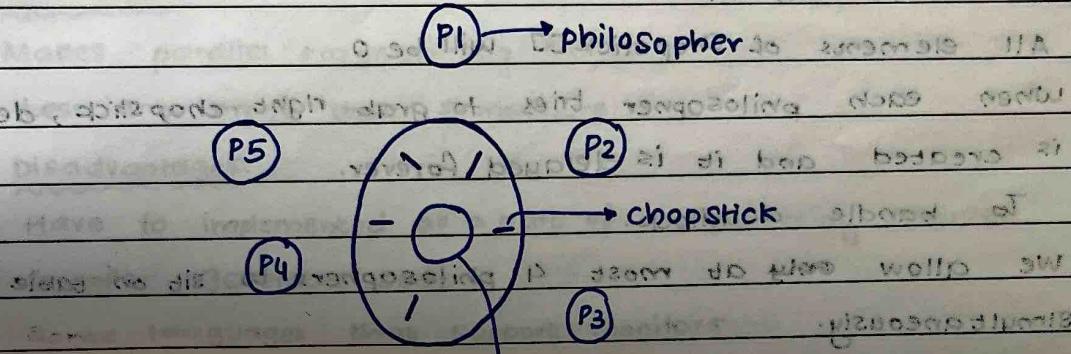
### 1 Chopstick between every neighbor.

- Picks up chopstick one at a time

- Needs both to eat, release both when done.

Dining Philosopher is a classic synchronisation problem

- Example of large class of concurrency control problems
- Represents the need to allocate several resources among several processes in a deadlock free and starvation free way.



- 5 philosophers share a common circular table. Bowl of noodles at the centre of table and 5 single chopsticks.

Shared dataset: Bowl of rice (dataset)

Semaphore - Chopstick [5] initialised to 1. (Binary Semaphore)

- Execute wait() operation if philosopher takes a chopstick
- Execute signal() operation if chopstick is released.

### \* Semaphore Solution

```
do {
```

    Wait C chopstick [i];      // philosopher i has picked up both chopsticks

    Wait C chopstick [(i+1) % 5];      // and eating function is

    // eat

    signal C chopstick [i];      // performed.

    signal C chopstick [(i+1) % 5];

    Once philosopher i is

    // think

    done eating, chopsticks

} while (1)

    are put down and it goes back in thinking state.

	M	T	W	T	F	S
MONDAY						
	Page No.:					
	Date:					

Page No.:	6
Date:	10/10/2023

- The algorithm guarantees that no two neighbors are eating simultaneously.

- But it is still rejected because of deadlock.

- If all 5 philosophers simultaneously decide to eat and each grabs left chopstick

- All elements of Chopstick[ ] will be 0

- When each philosopher tries to grab right chopstick, deadlock is created and it is delayed forever.

- To handle deadlock:

(1) We allow only at most 4 philosophers to sit on table simultaneously.

(2) Allow a philosopher to pick a fork only if both are available.

That is put the waiting operations in CS.

(3) Use an asymmetric solution.

Odd numbered philosophers will first pick up left chopstick and then right chopstick.

Even numbered philosopher will first pick up right chopstick and then left chopstick.

- These solutions still cannot get rid of starvation.

## \* Monitors:

= High-level synchronization construct

Programmer-defined operators

Declaration of variables

- Bodies of procedures

- monitors monitor-name

shared-p1 : public slots : { }

// shared variable declarations

Procedure P1 (...) { ... }

Procedure P2 (...) { ... }

:

Procedure P3 (...) { ... }

Initialization code (...) { ... }

3

- Collection of condition variables and procedures combined.
- Monitor ensures that only one process maybe active at a time.
- Programmer does not need to synchronize the code using constraint.
- $\alpha$ .Wait() and  $\alpha$ .Signal() conditions can be used.

### Advantages:

(i) Makes parallel programming easier

(ii) Less error prone than semaphore

### Disadvantages:

(i) Have to implemented as a part of programming language.

(ii) Compiler generates code

(iii) Some languages that support monitors

(iv) Compiler should know what OS is and other burdens

### Monitor Solution to dining Philosophers:

- Philosopher can be in 3 states :

enum {Thinking, Hungry, Eating} state[5];

- Philosopher will eat only if both chopsticks are available.

Philosopher i can be in variable state[i] = Eating only if two adjacent neighbors are not eating

$(\text{state}[(i+4)\%5] \neq \text{Eating}) \text{ and } (\text{state}[(i+1)\%5] \neq \text{Eating})$

- We also need to declare - condition self [5];

- ith philosopher can delay herself when she is hungry but unable to obtain chop sticks she needs.

- Distribution of chopsticks is maintained using monitor dining Philosophers

(See code from ppt)

S	C	R	T	W	T	F
A	V	U	O	M		
Date:						
Page No.:						

M	I	W	T	F
Page No.:				

## \* Deadlock

A process is in deadlock where every process is dependent on the other processes for an event to happen.

Deadlocks can arise because of 4 conditions:

### 1. Mutual Exclusion:

Only one process can use a resource at a time.

If another process requests the resource, the requesting must be delayed until the process is released.

### 2. Hold and Wait:

Process holding one resource wants to acquire one more resource that is held by another process.

### 3. No preemption:

Process cannot be preempted.

A resource can be released only by the process after process has completed its task.

### 4. Circular Wait:

There are  $n$  sets  $\{P_0, P_1, P_2, P_3, P_4, \dots, P_n\}$  of waiting processes such that  $P_0 \rightarrow P_1 \rightarrow P_2 \rightarrow \dots \rightarrow P_n \rightarrow P_0$

$P_0$  is waiting for resource held by  $P_1$ .

$P_1$  is waiting for resource held by  $P_2$ .

$P_2$  is waiting for resource held by  $P_3$ .

$P_3$  is waiting for resource held by  $P_4$ .

$P_4$  is waiting for resource held by  $P_0$ .

## \* Resource Allocation graph

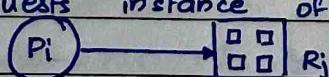
Request edge:  $P_i \rightarrow R_j$

Assignment edge:  $R_j \rightarrow P_i$

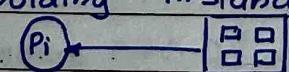
Resource type with 4 instances :



$P_i$  requests instance of  $R_j$ :



$P_i$  is holding instance of  $R_j$ :



- If graph contains no cycles then no deadlock.
- If graph contains a cycle: deadlock exists.
- If only one instance per resource then deadlock.
- If several instances per resource type then deadlock.

#### \* Deadlock Prevention:

- Avoid atleast one of the conditions causing deadlock.
- i. Mutual Exclusion: Mutual exclusion not required for shareable resources.
- eg: If there are several processes requesting to access a read-only file they can be granted simultaneous access.
- Mutual exclusion must hold for non-shareable resources.
- eg: A printer cannot be shared simultaneously.

#### 2. Hold and Wait:

- Must ensure that whenever a process makes a request for a resource no other process is holding it.
- Method 1:

  - Process must request and be allocated all the resources before execution begins.

- Method 2:

  - Allow process to request resources only when there is no one allocated to it.
  - Process can request additional resources, but it must release <sup>all</sup> the resources that are currently allocated.

eg: Consider a process that copies data from DVD drive to a file on disk, sorts the file and prints the result to a printer.

#### Method 1:

- All resources are requested before execution.
- Process requests DVD drive, file and printer initially.
- Printer is held for entire execution even though it was needed at the end.

### Method 2:

- Process requests only DVD and file initially.
- It copies the data from DVD to file and sorts it and immediately releases both of them.
- Process again requests for file and printer.
- After copying data to file and sorting it, it sends it to the printer and releases two resources again.
- For first method, resource utilization is low.
- For 2nd method, starvation is possible.

### No Preemption:

- Method 1:
  - If a process is holding some resources and requests another resource that cannot be allocated immediately then, all the resources being held currently must be released.
  - Process will restart only when its old resources and new ones are available on request.
- Method 2:
  - If a process requests some resource and if they are not available, we check if it is allocated to other resources that is itself waiting for some other resource.
  - If so, we preempt the desired resource from the waiting process and allocate it to requesting process.
  - If resources are neither available, nor being held, the requesting process must wait.
  - While waiting some of its resource get preempted only if they are requested.

### 4. Circular Wait:

- We impose a proper ordering of resource allocation.
- $P_1$  and  $P_2$  both need  $R_1$  and  $R_2$ . Locking both resources should be in a order like first lock  $R_1$  and then  $R_2$ . Whichever process gets lock  $R_1$  first locks  $R_2$ , and other process will wait for lock to release.

## \* Deadlock Avoidance

### \* Banker's Algorithm:

- A resource allocation and deadlock avoidance algorithm.
- It is named so to check whether loan can be sanctioned to a person or no.
- Suppose there are  $n$  account holders and total sum is  $S$ .
- If a money person applies for a loan, the bank will first subtract loan money from total amount. If the remaining  $S$  is greater than  $s$  then the loan is sanctioned.
- It is done because if all account holders came to withdraw money it could do it easily.
- Bank would never allocate money if it cannot satisfy the needs of all its customers.
- Bank always wants to be in safe state.

eg:	Process	Allocation	Max	Need	Available
	P <sub>0</sub>	a b c d	a b c d	a b c d	a b c d
	P <sub>1</sub>	1 0 0 0	1 5 0 0	1 5 0 0	1 5 2 0
	P <sub>2</sub>	1 3 5 4	2 3 5 6	1 0 0 2	
	P <sub>3</sub>	0 6 3 2	0 6 5 2	0 0 2 0	
	P <sub>4</sub>	0 0 1 4	0 6 5 6	0 6 4 2	

- Need = Max - Allocation

- Step 1: P<sub>0</sub>

Work = Allocation = [1, 5, 2, 0]

Need  $\leq$  Work, so safe

Sequence = <P<sub>0</sub>>

Work = Work + Allocation

= [1, 5, 3, 2]

T	F	F	F	F
---	---	---	---	---

Step 2: For P<sub>1</sub>

Work = [1, 5, 3, 2]

Need  $\not\leq$  Work, so unsafe

P<sub>1</sub> must wait.

Step 3: For P<sub>2</sub>

Work = [1, 5, 3, 2]

Sequence:

Need  $\leq$  Work i.e. <P<sub>0</sub>, P<sub>2</sub>>

Work = Work + Allocation

= [2, 8, 8, 6]

T	F	T	F	F
---	---	---	---	---

MON	TUE	WED	THU	FRI	SAT	SUN
Page No.:						
Date:						

M	T	W	T	F	S	S
Page No.:						
Date:						

For P<sub>3</sub>

$$Work = [2, 8, 8, 6]$$

Need ≤ Work, P<sub>3</sub> is safe

$$\text{Sequence} = \langle P_0, P_2, P_3 \rangle$$

$$Work = Work + Allocation$$

$$= [2, 14, 11, 8]$$

$$\text{Finish} = \begin{array}{|c|c|c|c|c|c|} \hline & True & False & True & True & False \\ \hline \end{array} \quad \begin{array}{|c|c|c|c|c|} \hline T & F & T & T & T \\ \hline \end{array}$$

For P<sub>4</sub>

$$Work = [2, 14, 11, 8]$$

Need ≤ Work, P<sub>4</sub> is safe

$$\text{Sequence} = \langle P_0, P_2, P_3, P_4 \rangle$$

$$Work = Work + Allocation$$

$$= [2, 14, 12, 12]$$

For P<sub>i</sub>

Work Need ≤ Work, P<sub>i</sub> is safe

$$\text{Sequence} = \langle P_0, P_2, P_3, P_4, P_i \rangle$$

$$Work = Work + Allocation$$

$$= [3, 4, 12, 12]$$

### \* Deadlock Detection:

- Allows system to enter deadlock state
- Single instance of each resource type: Wait-for-graph

Approaches: Remove all resource nodes and collapse appropriate edges with approaches. If a circle still exists in graph then deadlock present.

eg: see pdf

- Multiple instance: Banker's Algorithm

### \* Deadlock Recovery:

#### 1. Process Termination:

Abort one or more processes to break the circular wait.

• Abort all processes

• Abort one process at a time till DL is eliminated.

#### 2. Resource Preemption:

Selectively preempt some resources until deadlock cycle is broken.

Issues: starvation

Selecting a victim which resources and processes must be preempted, determine order & cost

Rollback return to safe state and restart for that state.

Starvation